

## ▼ Problem 2

### ▼ (a) (c) (d)

```
from mnist_tools import *
import numpy as np
import matplotlib.pyplot as plt
import time
from scipy.special import logsumexp

"""
Sigmoid function that takes a numpy array of any shape.
"""
def f(t) :
    return 1/(1+np.exp(-t))

"""
Forecast function which given the learned parameter vectors w and
the data x produces the forecasts.
"""
def h(w,x) :
    return f(np.dot(x,w))>0.5

"""
Computes the loss function L.
Parameters:
w: numpy array of length m containing the parameter vector
X: numpy array of shape (n,m) containing n data samples as rows (each row is a data
y: numpy array of length n containing the labels (0 or 1)
Returns:
A single float, the loss evaluated on the given arguments.
"""
def L(w,X,y):
    n = len(y)
    sum = 0
    for i in range(n):
        g_i = np.dot(w, X[i])
        term = y[i] * np.logaddexp(0, - g_i) + (1 - y[i]) * np.logaddexp(0, g_
        sum = sum + term / n
    return sum

"""
Tests the L function
"""
def test_L() :
    np.random.seed(1000)
    v = np.array([1000])
    w = np.random.randn(10)
    X = np.random.randn(20, 10)
    y = np.random.randint(0, 2, 20)
```

```

L1 = L(v, v, np.array([0]))
L2 = L(v, v, np.array([1]))
L3 = L(w, X, y)
assert np.abs(L1-1000000) < 1e-9
assert np.abs(L2) < 1e-9
assert np.abs(L3-1.08007365415) < 1e-9

```

"""

Computes the gradient of the loss function.

Parameters:

w: numpy array of length m containing the parameter vector

X: numpy array of shape (n,m) containing n data samples as rows (each row is a data

y: numpy array of length n containing the labels (0 or 1)

Returns:

A numpy vector of length m containing the gradient of the loss evaluated on the given arguments.

"""

```

def dL(w, X, y) :
    n = len(y)
    sum = 0
    for i in range(n):
        g_i = np.dot(w, X[i])
        term = X[i] * (f(g_i) - y[i])
        sum = sum + term / n
    return sum

```

"""

Tests the dL function

"""

```

def test_dL() :
    np.random.seed(1000)
    v = np.array([1000])
    w = np.random.randn(3)
    X = np.random.randn(200, 3)
    y = np.random.randint(0, 2, 200)
    dL1 = dL(v, v, np.array([0]))
    dL2 = dL(v, v, np.array([1]))
    dL3 = dL(w, X, y)
    assert np.abs(dL1-1000) < 1e-9
    assert np.abs(dL2) < 1e-9
    assert np.linalg.norm(dL3-np.array([-0.12669153, -0.00341384, 0.02274541])) < 1e-6

```

"""

Runs (batch) gradient descent with a backtracking line search to minimize L.

While typically this would include conditions/tolerances for how to stop the algorithm, here we only required a simplified implementation that has a given fixed number of steps.

Parameters:

w0: numpy array of length m containing the initial value of w

X: numpy array of shape (n,m) containing the n data samples as rows

y: numpy array of length n containing the labels (0 or 1)

num\_steps: number of gradient descent steps to run

alpha: Armijo constant used to make sure the L function sufficiently decreases on each iteration

beta: backtracking line search constant that determines how much to shrink the step

size parameter by each time

Returns: the tuple w,ws where

w: numpy array of length m containing the final value of w

ws: a python list of num\_steps numpy arrays of length m containing the w-values computed at each iteration

"""

```
def gradient_descent(w0, X, y, num_steps=200, alpha=0.01, beta=0.5) :
    w_s = []
    w = w0
    for i in range(num_steps):
        t = 1
        while (L(w, X, y) - L(w - t * dL(w, X, y), X, y) - alpha * t * n
               t = t * beta
        print("t = ", t)
        w = w - t * dL(w, X, y)
        w_s.append(w)
    return w, w_s
```

"""

Standardizes the training and test data using the training data to compute the mean and standard deviation.

"""

```
def standardize(train, test) :
    m = np.mean(train, axis=0)
    std = np.std(train, axis=0)
    std[np.abs(std)<1e-9] = 1
    return (train-m)/std, (test-m)/std, m, std
```

"""

Runs the optimization and creates the plots

"""

```
def run(name, fun, train_x, train_y, test_x, test_y, mean, std) :
    t = time.time()
    g_w, g_ws = fun(np.zeros(train_x.shape[1]), train_x, train_y)
    print('%s Time = %fs'%(name, time.time()-t))
    print('%s Training Loss = %f'%(name, L(g_w, train_x, train_y)))
    test_err = np.sum(np.abs(h(g_w, test_x)-test_y))*1.0/test_x.shape[0]
    print('%s Test Error = %f'%(name, test_err))
    ls = [L(w, train_x, train_y) for w in g_ws]
    tls = [L(w, test_x, test_y) for w in g_ws]
    terr = [np.sum(np.abs(h(w, test_x)-test_y))/test_x.shape[0] for w in g_ws]
    plt.plot(ls)
    plt.title('%s Training Loss'%name)
    plt.savefig('%s_Train_Loss.pdf'%name, bbox_inches='tight')
    plt.close()
    plt.title('%s Test Loss'%name)
    plt.plot(tls)
    plt.savefig('%s_Test_Loss.pdf'%name, bbox_inches='tight')
    plt.close()
    plt.plot(terr)
```

```

plt.title('%s Test Error'%name)
plt.savefig('%s_Test_Error.pdf'%name, bbox_inches='tight')
plt.close()

def main() :
    test_L()
    test_dL()

    train = load_train_data("mnist_all.mat")
    test = load_test_data("mnist_all.mat")
    print('Using %d training examples and %d test examples'%(train.shape[0], test.shape[0])
    #We will determine if the image is a '5' or not
    train[:, -1] = train[:, -1] == 5
    test[:, -1] = test[:, -1] == 5
    train_x, train_y = train[:, :-1], train[:, -1]
    test_x, test_y = test[:, :-1], test[:, -1]
    train_x, test_x, mean, std = standardize(train_x, test_x)

    run('GD', gradient_descent, train_x, train_y, test_x, test_y, mean, std)

if __name__ == "__main__" :
    main()

```